

# Homework 2: Food Safety (50 Pts)

## Cleaning and Exploring Data with Pandas

### This Assignment

In this homework, we will investigate restaurant food safety scores for restaurants in San Francisco. The scores and violation information have been [made available by the San Francisco Department of Public Health](#). The main goal for this assignment is to walk through the process of Data Cleaning and EDA.

As we clean and explore these data, you will gain practice with:

- Reading simple csv files and using Pandas
- Working with data at different levels of granularity
- Identifying the type of data collected, missing values, anomalies, etc.
- Exploring characteristics and distributions of individual variables

```
In [ ]: import numpy as np
import pandas as pd

import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
plt.style.use('fivethirtyeight')

import zipfile
from pathlib import Path
import os # Used to interact with the file system
```

### Importing and Verifying Data

There are several tables in the data folder. Let's attempt to load `bus.csv`, `ins2vio.csv`, `ins.csv`, and `vio.csv` into pandas dataframes with the following names: `bus`, `ins2vio`, `ins`, and `vio` respectively.

*Note:* Because of character encoding issues one of the files (`bus`) will require an additional argument `encoding='ISO-8859-1'` when calling `pd.read_csv`.

```
In [ ]: # path to directory containing data
dsDir = Path('data')

bus = pd.read_csv(dsDir/'bus.csv', encoding='ISO-8859-1')
ins2vio = pd.read_csv(dsDir/'ins2vio.csv')
ins = pd.read_csv(dsDir/'ins.csv')
vio = pd.read_csv(dsDir/'vio.csv')
```

Now that you've read in the files, let's try some `pd.DataFrame` methods ([docs](#)). Use the `DataFrame.head` method to show the top few lines of the `bus`, `ins`, and `vio` dataframes. To show multiple return outputs in one single cell, you can use `display()`. Currently, running the cell below will display the first few lines of the `bus` dataframe.

```
In [ ]: bus.head()
```

```
Out[ ]:
```

	business id column	name	address	city	state	postal_code	lati
0	1000	HEUNG YUEN RESTAURANT	3279 22nd St	San Francisco	CA	94110	37.75
1	100010	ILLY CAFFE SF_PIER 39	PIER 39 K-106-B	San Francisco	CA	94133	-9999.00
2	100017	AMICI'S EAST COAST PIZZERIA	475 06th St	San Francisco	CA	94103	-9999.00
3	100026	LOCAL CATERING	1566 CARROLL AVE	San Francisco	CA	94124	-9999.00
4	100030	OUI OUI! MACARON	2200 JERROLD	San Francisco	CA	94124	-9999.00

The `DataFrame.describe` method can also be handy for computing summaries of numeric columns of our dataframes. Try it out with each of our 4 dataframes. Below, we have used the method to give a summary of the `bus` dataframe.

```
In [ ]: bus.describe()
```

```
Out[ ]:
```

	business id column	latitude	longitude	phone_number
<b>count</b>	6253.000000	6253.000000	6253.000000	6.253000e+03
<b>mean</b>	60448.948984	-5575.337966	-5645.817699	4.701819e+09
<b>std</b>	36480.132445	4983.390142	4903.993683	6.667508e+09
<b>min</b>	19.000000	-9999.000000	-9999.000000	-9.999000e+03
<b>25%</b>	18399.000000	-9999.000000	-9999.000000	-9.999000e+03
<b>50%</b>	75685.000000	-9999.000000	-9999.000000	-9.999000e+03
<b>75%</b>	90886.000000	37.776494	-122.421553	1.415533e+10
<b>max</b>	102705.000000	37.824494	0.000000	1.415988e+10

Now, we perform some sanity checks for you to verify that the data was loaded with the correct structure. Run the following cells to load some basic utilities (you do not need to change these at all):

First, we check the basic structure of the data frames you created:

```
In [ ]: assert all(bus.columns == ['business id column', 'name', '
                                             'latitude', 'longitude', 'phone
assert 6250 <= len(bus) <= 6260

assert all(ins.columns == ['iid', 'date', 'score', 'type'])
assert 26660 <= len(ins) <= 26670

assert all(vio.columns == ['description', 'risk_category',
assert 60 <= len(vio) <= 65

assert all(ins2vio.columns == ['iid', 'vid'])
assert 40210 <= len(ins2vio) <= 40220
```

Next we'll check that the statistics match what we expect. The following are hard-coded statistical summaries of the correct data.

```
In [ ]: bus_summary = pd.DataFrame(**{'columns': ['business id col',
'data': {'business id column': {'50%': 75685.0, 'max': 102705.0, 'min': 19.0},
'latitude': {'50%': -9999.0, 'max': 37.824494, 'min': -9999.0},
'longitude': {'50%': -9999.0, 'max': 0.0, 'min': -9999.0}},
'index': ['min', '50%', 'max']})

ins_summary = pd.DataFrame(**{'columns': ['score'],
'data': {'score': {'50%': 76.0, 'max': 100.0, 'min': -1.0},
'index': ['min', '50%', 'max']})

vio_summary = pd.DataFrame(**{'columns': ['vid'],
'data': {'vid': {'50%': 103135.0, 'max': 103177.0, 'min': 103102.0},
'index': ['min', '50%', 'max']})

from IPython.display import display

print('What we expect from your Businesses dataframe:')
display(bus_summary)
print('What we expect from your Inspections dataframe:')
display(ins_summary)
print('What we expect from your Violations dataframe:')
display(vio_summary)
```

What we expect from your Businesses dataframe:

	business id column	latitude	longitude
min	19.0	-9999.000000	-9999.0
50%	75685.0	-9999.000000	-9999.0
max	102705.0	37.824494	0.0

What we expect from your Inspections dataframe:

	score
min	-1.0
50%	76.0
max	100.0

What we expect from your Violations dataframe:

	vid
min	103102.0

50% 103135.0

max 103177.0

The code below defines a testing function that we'll use to verify that your data has the same statistics as what we expect. Run these cells to define the function. The `df_allclose` function has this name because we are verifying that all of the statistics for your dataframe are close to the expected values. Why not `df_allegal`? It's a bad idea in almost all cases to compare two floating point values like 37.780435, as rounding error can cause spurious failures.

```
In [ ]: """Run this cell to load this utility comparison function
tests below

Do not modify the function in any way.
"""

def df_allclose(actual, desired, columns=None, rtol=5e-2):
    """Compare selected columns of two dataframes on a few
    Compute the min, median and max of the two dataframes
    that they match numerically to the given relative tolerance.

    If they don't match, an AssertionError is raised (by default).
    """
    # summary statistics to compare on
    stats = ['min', '50%', 'max']

    # For the desired values, we can provide a full DF with
    # the actual data, or pre-computed summary statistics.
    # We assume a pre-computed summary was provided if columns
    # `desired` *must* have the same structure as the actual
    if columns is None:
        des = desired
        columns = desired.columns
    else:
        des = desired[columns].describe().loc[stats]

    # Extract summary stats from actual DF
    act = actual[columns].describe().loc[stats]

    return np.allclose(act, des, rtol)
```

## Question 1a: Identifying Issues with the Data

Use the `head` command on your three files again. This time, describe at least one potential problem with the data you see. Consider issues with missing values and bad data.

```
In [ ]: bus.head()
```

```
Out[ ]:
```

	business id column	name	address	city	state	postal_code	lati
0	1000	HEUNG YUEN RESTAURANT	3279 22nd St	San Francisco	CA	94110	37.75
1	100010	ILLY CAFFE SF_PIER 39	PIER 39 K-106-B	San Francisco	CA	94133	-9999.00
2	100017	AMICI'S EAST COAST PIZZERIA	475 06th St	San Francisco	CA	94103	-9999.00
3	100026	LOCAL CATERING	1566 CARROLL AVE	San Francisco	CA	94124	-9999.00
4	100030	OUI OUI! MACARON	2200 JERROLD AVE STE C	San Francisco	CA	94124	-9999.00

```
In [ ]: ins.head()
```

```
Out[ ]:
```

	iid	date	score	type
0	100010_20190329	03/29/2019 12:00:00 AM	-1	New Construction

<b>1</b>	100010_20190403	04/03/2019 12:00:00 AM	100	Routine - Unscheduled
<b>2</b>	100017_20190417	04/17/2019 12:00:00 AM	-1	New Ownership
<b>3</b>	100017_20190816	08/16/2019 12:00:00 AM	91	Routine - Unscheduled
<b>4</b>	100017_20190826	08/26/2019 12:00:00 AM	-1	Reinspection/Followup

```
In [ ]: vio.head()
```

		<b>description</b>	<b>risk_category</b>	<b>vid</b>
<b>0</b>	Consumer advisory not provided for raw or unde...		Moderate Risk	103128
<b>1</b>		Contaminated or adulterated food	High Risk	103108
<b>2</b>	Discharge from employee nose mouth or eye		Moderate Risk	103117
<b>3</b>		Employee eating or smoking	Moderate Risk	103118
<b>4</b>		Food in poor condition	Moderate Risk	103123

**ANSWER:** The phone number in the first row of bus is -9999 which is unreasonable. Also, the latitude and longitude in this file is also weird which might cause issues when doing analysis about these values.

We will explore each file in turn, including determining its granularity and primary keys and exploring many of the variables individually. Let's begin with the businesses file, which has been read into the `bus` dataframe.

---

# Question 1b: Examining the Business Data File

From its name alone, we expect the `bus.csv` file to contain information about the restaurants. Let's investigate the granularity of this dataset.

```
In [ ]: bus.head()
```

```
Out[ ]:
```

	business id column	name	address	city	state	postal_code	lati
0	1000	HEUNG YUEN RESTAURANT	3279 22nd St	San Francisco	CA	94110	37.75
1	100010	ILLY CAFFE SF_PIER 39	PIER 39 K-106-B	San Francisco	CA	94133	-9999.00
2	100017	AMICI'S EAST COAST PIZZERIA	475 06th St	San Francisco	CA	94103	-9999.00
3	100026	LOCAL CATERING	1566 CARROLL AVE	San Francisco	CA	94124	-9999.00
4	100030	OUI OUI! MACARON	2200 JERROLD AVE STE C	San Francisco	CA	94124	-9999.00

The `bus` dataframe contains a column called `business id column` which probably corresponds to a unique business id. However, we will first rename that column to `bid` for simplicity.

```
In [ ]: bus = bus.rename(columns={"business id column": "bid"})
```

Examining the entries in `bus`, is the `bid` unique for each record (i.e. each row of data)? Your code should compute the answer, i.e. don't just hard code `True` or `False`.

Hint: use `value_counts()` or `unique()` to determine if the `bid` series has any duplicates.



```
In [ ]: is_bid_unique = bus.shape[0] == bus['bid'].unique().size
```

## Question 1c

We will now work with some important fields in `bus`. In the two cells below create the following **two numpy arrays**:

1. Assign `top_names` to the top 5 most frequently used business names, from most frequent to least frequent.
2. Assign `top_addresses` to the top 5 addresses where businesses are located, from most popular to least popular.

Hint: you may find `value_counts()` helpful.

## Step 1

```
In [ ]: top_names = bus['name'].value_counts().sort_values(ascending=False)
top_addresses = bus['address'].value_counts().sort_values(ascending=False)
top_names, top_addresses
```

```
Out[ ]: (Peet's Coffee & Tea      20
          Starbucks Coffee      13
          Jamba Juice           10
          McDonald's            10
          STARBUCKS             9
          Name: name, dtype: int64,
          Off The Grid          39
          428 11th St           34
          3251 20th Ave         17
          2948 Folsom St        17
          Pier 41                16
          Name: address, dtype: int64)
```

# Question 1d

Based on the above exploration, answer each of the following questions about `bus` by assigning your answers to the corresponding variables

1. What does each record represent?
2. What is the minimal primary key?

Minimal primary key: a minimal set of attributes (columns) that uniquely specify a row in a table

```
In [ ]: # What does each record represent? Valid answers are:
#       "One location of a restaurant."
#       "A chain of restaurants."
#       "A city block."
q1d_part1 = "One location of a restaurant."

# What is the minimal primary key? Valid answers are:
#       "bid"
#       "bid, name"
#       "bid, name, address"
q1d_part2 = "bid"
```

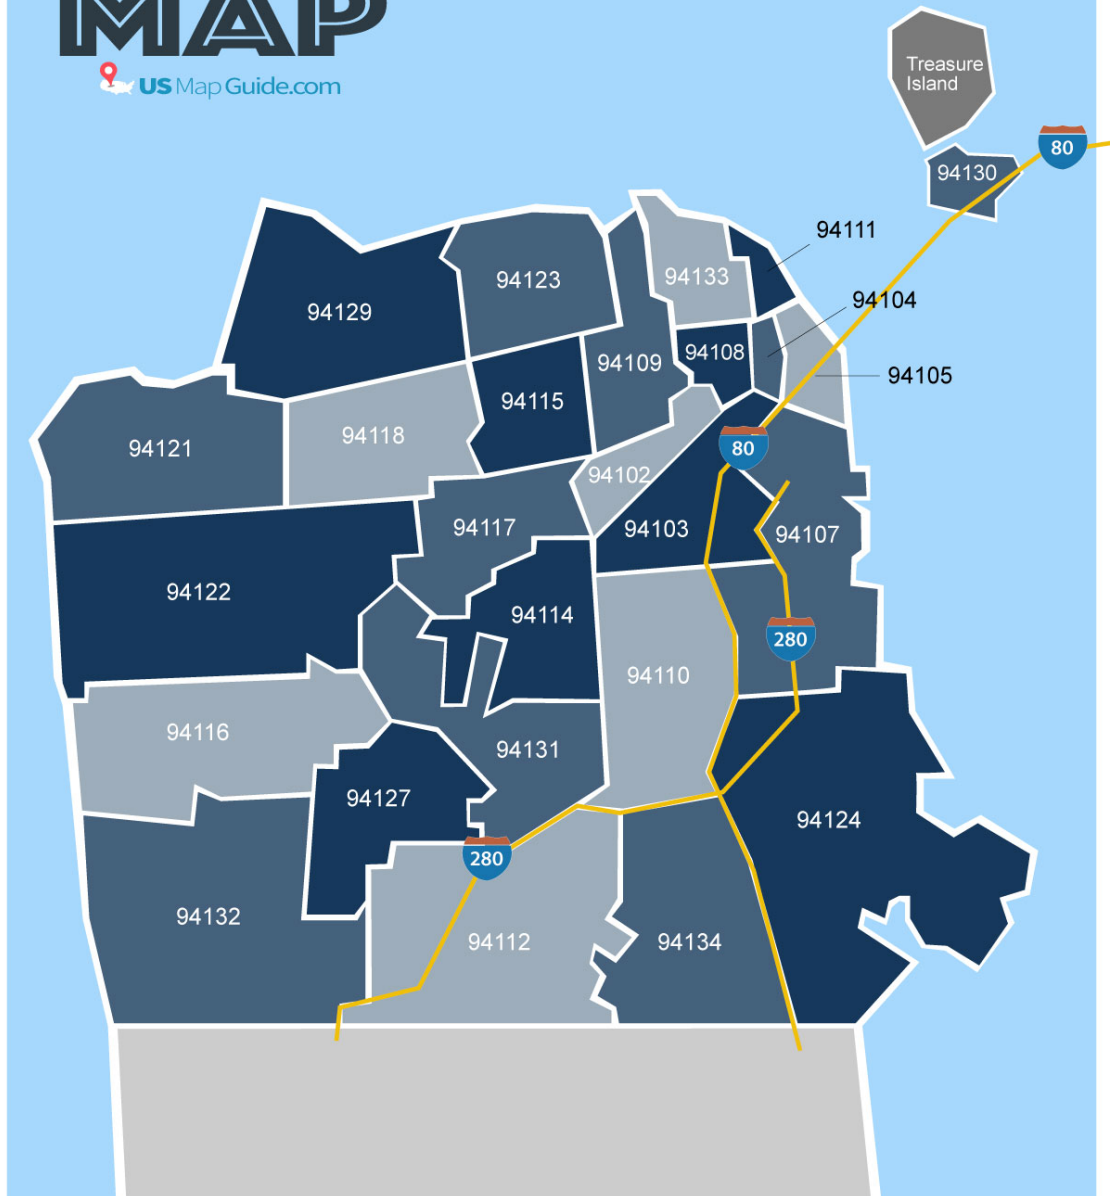
## 2: Cleaning the Business Data Postal Codes

The business data contains postal code information that we can use to aggregate the ratings over regions of the city. Let's examine and clean the postal code field. The postal code (sometimes also called a ZIP code) partitions the city into regions:

SAN FRANCISCO

# Printable MAP

 USMapGuide.com



## Question 2a

How many restaurants are in each ZIP code?

In the cell below, create a **series** where the index is the postal code and the value is the number of records with that postal code in

descending order of count. You may need to use `groupby()`, `size()`, or `value_counts()`. Do you notice any odd/invalid zip codes?

```
In [ ]: zip_counts = bus.groupby('postal_code').size().sort_values
print(zip_counts.to_string())
```

```
postal_code
94103      562
94110      555
94102      456
94107      408
94133      398
94109      382
94111      259
94122      255
94105      249
94118      231
94115      230
94108      229
94124      218
94114      200
-9999      194
94112      192
94117      189
94123      177
94121      157
94104      142
94132      132
94116       97
94158       90
94134       82
94127       67
94131       49
94130        8
94143        5
94013        2
94188        2
CA          2
94301        2
94101        2
95122        1
941033148    1
95133        1
95132        1
94102-5917    1
94014        1
941         1
94080        1
94105-2907    1
```

92672	1
64110	1
00000	1
94105-1420	1
941102019	1
95117	1
95112	1
95109	1
95105	1
94901	1
94621	1
94602	1
94544	1
94518	1
94117-3504	1
94120	1
94122-1909	1
94123-3106	1
94124-1917	1
94129	1
Ca	1

## Question 2b

Answer the question about the `postal_code` column in the `bus` dataframe.

1. What Python data type is used to represent a ZIP code?

*Note:* ZIP codes and postal codes are the same thing.

Please write your answers in the variables below:

```
In [ ]: # What Python data type is used to represent a ZIP code?
#       "str"
#       "int"
#       "bool"
#       "float"
q2b = "str"
```

---

## Question 2c

In question 2a we noticed a large number of potentially invalid ZIP codes (e.g., "Ca"). These are likely due to data entry errors. To get a better understanding of the potential errors in the zip codes we will:

1. Import a list of valid San Francisco ZIP codes by using `pd.read_json` to load the file `data/sf_zipcodes.json` and extract a **series** of type `str` containing the valid ZIP codes. *Hint: set `dtype` when invoking `read_json`.*
2. Construct a `DataFrame` containing only the businesses which DO NOT have valid ZIP codes. You will probably want to use the `Series.isin` function.

## Step 1

```
In [ ]: valid_zips = pd.read_json("data/sf_zipcodes.json", dtype="str")
valid_zips.head()
```

```
Out[ ]: zip_codes
```

0	94102
1	94103
2	94104
3	94105
4	94107

## Step 2

```
In [ ]: invalid_zip_bus = bus[~bus["postal_code"].isin(valid_zips)]
invalid_zip_bus.head(20)
```

```
Out[ ]:
```

	bid	name	address	city	state	postal_code	
22	100126	Lamas Peruvian Food Truck	Private Location	San Francisco	CA	-9999	-99
68	100417	COMPASS ONE, LLC	1 MARKET ST. FL	San Francisco	CA	94105-1420	-99
96	100660	TEAPENTER	1518 IRVING ST	San Francisco	CA	94122-1909	-99
109	100781	LE CAFE DU SOLEIL	200 FILLMORE ST	San Francisco	CA	94117-3504	-99

<b>144</b>	101084	Deli North 200	1 Warriors Way Level 300 North East	San Francisco	CA	94518	-99
<b>156</b>	101129	Vendor Room 200	1 Warriors Way Level 300 South West	San Francisco	CA	-9999	-99
<b>177</b>	101192	Cochinita #2	2 Marina Blvd Fort Mason	San Francisco	CA	-9999	-99
<b>276</b>	102014	DROPBOX (Section 3, Floor 7)	1800 Owens St	San Francisco	CA	-9999	-99
<b>295</b>	102245	Vessell CA Operations (#4)	2351 Mission St	San Francisco	CA	-9999	-99
<b>298</b>	10227	The Napper Tandy	3200 24th St	San Francisco	CA	-9999	
<b>320</b>	10372	BERNAL HEIGHTS NEIGHBORHOOD CENTER	515 CORTLAND AVE	San Francisco	CA	-9999	
<b>321</b>	10373	El Tonayense #1	1717 Harrison St	San Francisco	CA	-9999	
<b>322</b>	10376	Good Frikin Chicken	10 29th St	San Francisco	CA	-9999	
<b>324</b>	10406	Sunset Youth Services	3918 Judah St	San Francisco	CA	-9999	
<b>357</b>	11416	El Beach Burrito	3914 Judah St	San Francisco	CA	-9999	
<b>381</b>	12199	El Gallo Giro	3055 23rd St	San Francisco	CA	-9999	
<b>384</b>	12344	The Village Market & Pizza	750 Font Blvd	San Francisco	CA	-9999	
<b>406</b>	13062	Everett Middle School	450 Church St	San Francisco	CA	-9999	
<b>434</b>	13753	Taboun	203 Parnassus Ave	San Francisco	CA	-9999	
<b>548</b>	17423	Project Open	100	San	CA	-9999	

---

## Question 2d

In the previous question, many of the businesses had a common invalid postal code that was likely used to encode a MISSING postal code. Do they all share a potentially "interesting address"?

In the following cell, construct a **series** that counts the number of businesses at each `address` that have this single likely MISSING postal code value. Order the series in descending order by count.

After examining the output, please answer the following question (2e) by filling in the appropriate variable. If we were to drop businesses with MISSING postal code values would a particular class of business be affected? If you are unsure try to search the web for the most common addresses.

```
In [ ]: missing_zip_address_count = invalid_zip_bus.groupby(['addr  
missing_zip_address_count.head()
```

```
Out[ ]: address  
Off The Grid          39  
Off the Grid          10  
OTG                   4  
OFF THE GRID          4  
Approved Private Locations  3  
dtype: int64
```

---

## Question 2e



Examine the `invalid_zip_bus` dataframe we computed above and look at the businesses that DO NOT have the special MISSING ZIP code value. Some of the invalid postal codes are just the full 9 digit code rather than the first 5 digits. Create a new column named `postal5` in the original `bus` dataframe which contains only the first 5 digits of the `postal_code` column. Finally, for any of the `postal5` ZIP code entries that were not a valid San Francisco ZIP Code (according to `valid_zips`) set the entry to `None`.

```
In [ ]: bus['postal5'] = None
bus['postal5'] = bus['postal_code'].str[:5]
bus.loc[invalid_zip_bus.index, 'postal5'] = None

# Checking the corrected postal5 column
bus.loc[invalid_zip_bus.index, ['bid', 'name', 'postal_cod
```

```
Out[ ]:
```

	<b>bid</b>	<b>name</b>	<b>postal_code</b>	<b>postal5</b>
<b>22</b>	100126	Lamas Peruvian Food Truck	-9999	None
<b>68</b>	100417	COMPASS ONE, LLC	94105-1420	None
<b>96</b>	100660	TEAPENTER	94122-1909	None
<b>109</b>	100781	LE CAFE DU SOLEIL	94117-3504	None
<b>144</b>	101084	Deli North 200	94518	None
...	...	...	...	...
<b>6173</b>	99369	HOTEL BIRON	94102-5917	None
<b>6174</b>	99376	Mashallah Halal Food truck Ind	-9999	None
<b>6199</b>	99536	FAITH SANDWICH #2	94105-2907	None
<b>6204</b>	99681	Twister	95112	None
<b>6241</b>	99819	CHESTNUT DINER	94123-3106	None

230 rows × 4 columns

## 3: Investigate the Inspection Data

Let's now turn to the inspection DataFrame. Earlier, we found that `ins` has 4 columns named `iid`, `score`, `date` and `type`. In this section, we determine the granularity of `ins` and investigate the kinds of information provided for the inspections.

Let's start by looking again at the first 5 rows of `ins` to see what we're working with.

```
In [ ]: ins.head(5)
```

```
Out[ ]:
```

	iid	date	score	type
0	100010_20190329	03/29/2019 12:00:00 AM	-1	New Construction
1	100010_20190403	04/03/2019 12:00:00 AM	100	Routine - Unscheduled
2	100017_20190417	04/17/2019 12:00:00 AM	-1	New Ownership
3	100017_20190816	08/16/2019 12:00:00 AM	91	Routine - Unscheduled
4	100017_20190826	08/26/2019 12:00:00 AM	-1	Reinspection/Followup

---

### Question 3a

The column `iid` probably corresponds to an inspection id. Is it a primary key? Write an expression (line of code) that evaluates to `True` or `False` based on whether all the values are unique.

```
In [ ]: is_ins_iid_a_primary_key = ins.shape[0] == ins['iid'].unique().shape[0]
```

```
In [ ]: ins.shape[0]
```

```
Out[ ]: 26663
```

---

## Question 3b

The column `iid` appears to be the composition of two numbers and the first number looks like a business id.

**Part 1.:** Create a new column called `bid` in the `ins` dataframe containing just the business id. You will want to use `ins['iid'].str` operations to do this. Also be sure to convert the type of this column to `int`

**Part 2.:** Then compute how many values in this new column are invalid business ids (i.e. do not appear in the `bus['bid']` column). Consider using the `pd.Series.isin` function.

**No python for loops or list comprehensions required!**

### Part 1

```
In [ ]: ins['bid'] = ins['iid'].str.split('_').str.get(0).astype(int)
```

### Part 2

```
In [ ]: invalid_bid_count = ins[~ins['bid'].isin(bus['bid'])].shape[0]
```

---

## Question 3c

What if we are interested in a time component of the inspection data?  
We need to examine the date column of each inspection.

**Part 1:** What is the type of the individual `ins['date']` entries? You may want to grab the very first entry and use the `type` function in python.

**Part 2:** Use `pd.to_datetime` to create a new `ins['timestamp']` column containing of `pd.Timestamp` objects. These will allow us to do more date manipulation.

**Part 3:** What are the earliest and latest dates in our inspection data?  
*Hint: you can use `min` and `max` on dates of the correct type.*

**Part 4:** We probably want to examine the inspections by year. Create an additional `ins['year']` column containing just the year of the inspection. Consider using `pd.Series.dt.year` to do this.

**No python `for` loops or list comprehensions required!**

### Part 1

```
In [ ]: ins_date_type = type(ins.loc[0, 'date'])
ins_date_type
```

```
Out[ ]: str
```

### Part 2

```
In [ ]: ins['timestamp'] = pd.to_datetime(ins['date'])
```

### Part 3

```
In [ ]: earliest_date = ins['timestamp'].min()
latest_date = ins['timestamp'].max()

print("Earliest Date:", earliest_date)
print("Latest Date:", latest_date)
```

Earliest Date: 2016-10-04 00:00:00  
Latest Date: 2019-11-28 00:00:00

## Part 4

```
In [ ]: ins['year'] = ins['timestamp'].dt.year
```

```
In [ ]: ins.head()
```

```
Out[ ]:
```

	iid	date	score	type	bid	time
0	100010_20190329	03/29/2019 12:00:00 AM	-1	New Construction	100010	20
1	100010_20190403	04/03/2019 12:00:00 AM	100	Routine - Unscheduled	100010	20
2	100017_20190417	04/17/2019 12:00:00 AM	-1	New Ownership	100017	20
3	100017_20190816	08/16/2019 12:00:00 AM	91	Routine - Unscheduled	100017	20
4	100017_20190826	08/26/2019 12:00:00 AM	-1	Reinspection/Followup	100017	20

---

## Question 3d

What is the relationship between the type of inspection over the 2016 to 2019 timeframe?

### Part 1

Construct the following table by

1. Using the `pivot_table` containing the number ( `size` ) of inspections for the given `type` and `year` .
2. Adding an extra `Total` column to the result using `sum`
3. Sort the results in descending order by the `Total` .

year	2016	2017	2018	2019	Total
type					
<b>Routine - Unscheduled</b>	966	4057	4373	4681	14077
<b>Reinspection/Followup</b>	445	1767	1935	2292	6439
<b>New Ownership</b>	99	506	528	459	1592
<b>Complaint</b>	91	418	512	437	1458
<b>New Construction</b>	102	485	218	189	994
<b>Non-inspection site visit</b>	51	276	253	231	811
<b>New Ownership - Followup</b>	0	45	219	235	499
<b>Structural Inspection</b>	1	153	50	190	394
<b>Complaint Reinspection/Followup</b>	19	68	70	70	227
<b>Foodborne Illness Investigation</b>	1	29	50	35	115
<b>Routine - Scheduled</b>	0	9	8	29	46
<b>Administrative or Document Review</b>	2	1	1	0	4
<b>Multi-agency Investigation</b>	0	0	1	2	3
<b>Special Event</b>	0	3	0	0	3
<b>Community Health Assessment</b>	1	0	0	0	1

**No python `for` loops or list comprehensions required!**

```
In [ ]: ins_pivot = ins.pivot_table(index='type', columns='year',
ins_pivot['Total'] = ins_pivot.sum(axis='columns')

...
ins_pivot_sorted = ins_pivot.sort_values('Total', ascending=False)
ins_pivot_sorted
```

```
Out[ ]:                                     score  Total
type
year  2016  2017  2018  2019
```

<b>Routine - Unscheduled</b>	966	4057	4373	4681	14077
<b>Reinspection/Followup</b>	445	1767	1935	2292	6439
<b>New Ownership</b>	99	506	528	459	1592
<b>Complaint</b>	91	418	512	437	1458
<b>New Construction</b>	102	485	218	189	994
<b>Non-inspection site visit</b>	51	276	253	231	811
<b>New Ownership - Followup</b>	0	45	219	235	499
<b>Structural Inspection</b>	1	153	50	190	394
<b>Complaint Reinspection/Followup</b>	19	68	70	70	227
<b>Foodborne Illness Investigation</b>	1	29	50	35	115
<b>Routine - Scheduled</b>	0	9	8	29	46
<b>Administrative or Document Review</b>	2	1	1	0	4
<b>Multi-agency Investigation</b>	0	0	1	2	3
<b>Special Event</b>	0	3	0	0	3
<b>Community Health Assessment</b>	1	0	0	0	1

## Part 2

Based on the above analysis, which year appears to have had a lot of businesses in newly constructed buildings?

```
In [ ]: year_of_new_construction = 2017
```

## Question 3e

Let's examine the inspection scores `ins['score']`

```
In [ ]: ins['score'].value_counts().head()
```

```
Out[ ]: -1      12632
        100      1993
```

```
96      1681
92      1260
94      1250
```

```
Name: score, dtype: int64
```

There are a large number of inspections with the 'score' of -1 . These are probably missing values. Let's see what type of inspections have scores and which do not. Create the following dataframe using steps similar to the previous question, and assign it to the variable `ins_missing_score_pivot` .

You should observe that inspection scores appear only to be assigned to `Routine - Unscheduled` inspections.

Missing Score	False	True	Total
type			
Routine - Unscheduled	14031	46	14077
Reinspection/Followup	0	6439	6439
New Ownership	0	1592	1592
Complaint	0	1458	1458
New Construction	0	994	994
Non-inspection site visit	0	811	811
New Ownership - Followup	0	499	499
Structural Inspection	0	394	394
Complaint Reinspection/Followup	0	227	227
Foodborne Illness Investigation	0	115	115
Routine - Scheduled	0	46	46
Administrative or Document Review	0	4	4
Multi-agency Investigation	0	3	3
Special Event	0	3	3
Community Health Assessment	0	1	1

```
In [ ]: ins['Missing score'] = ins['score'] != -1
ins_missing_score_pivot = ins.pivot_table(index='type', co
ins_missing_score_pivot['Total'] = ins_missing_score_pivot
ins_missing_score_pivot.sort_values('Total', ascending=False)
```



Out[ ]:

	Missing score	False	True	Total
type				
<b>Routine - Unscheduled</b>	46	14031	14077	
<b>Reinspection/Followup</b>	6439	0	6439	
<b>New Ownership</b>	1592	0	1592	
<b>Complaint</b>	1458	0	1458	
<b>New Construction</b>	994	0	994	
<b>Non-inspection site visit</b>	811	0	811	
<b>New Ownership - Followup</b>	499	0	499	
<b>Structural Inspection</b>	394	0	394	
<b>Complaint Reinspection/Followup</b>	227	0	227	
<b>Foodborne Illness Investigation</b>	115	0	115	
<b>Routine - Scheduled</b>	46	0	46	
<b>Administrative or Document Review</b>	4	0	4	
<b>Multi-agency Investigation</b>	3	0	3	
<b>Special Event</b>	3	0	3	
<b>Community Health Assessment</b>	1	0	1	

Notice that inspection scores appear only to be assigned to **Routine - Unscheduled** inspections. It is reasonable that for inspection types such as **New Ownership** and **Complaint** to have no associated inspection scores, but we might be curious why there are no inspection scores for the **Reinspection/Followup** inspection type.

---

## 4: Joining Data Across Tables

In this question we will start to connect data across multiple tables. We will be using the `merge` function.

---

### Question 4a

Let's figure out which restaurants had the lowest scores. Before we proceed, let's filter out missing scores from `ins` so that negative scores don't influence our results.

```
In [ ]: ins = ins[ins["score"] > 0]
```

We'll start by creating a new dataframe called `ins_named`. It should be exactly the same as `ins`, except that it should have the name and address of every business, as determined by the `bus` dataframe. If a `business_id` in `ins` does not exist in `bus`, the name and address should be given as `NaN`.

*Hint:* Use the `merge` method to join the `ins` dataframe with the appropriate portion of the `bus` dataframe. See the official [documentation](#) on how to use `merge`.

*Note:* For quick reference, a pandas 'left' join keeps the keys from the left frame, so if `ins` is the left frame, all the keys from `ins` are kept and if a set of these keys don't have matches in the other frame, the columns from the other frame for these "unmatched" key rows contains NaNs.

```
In [ ]: ins_named = pd.merge(ins, bus[['bid', 'name', 'address']],
                             ins_named.head())
```

```
Out[ ]:
```

iid	date	score	type	bid	timestamp
-----	------	-------	------	-----	-----------

---

0	100010_20190403	04/03/2019 12:00:00 AM	100	Routine - Unscheduled	100010	2019-04-03	2
1	100017_20190816	08/16/2019 12:00:00 AM	91	Routine - Unscheduled	100017	2019-08-16	2
2	100041_20190520	05/20/2019 12:00:00 AM	83	Routine - Unscheduled	100041	2019-05-20	2
3	100055_20190425	04/25/2019 12:00:00 AM	98	Routine - Unscheduled	100055	2019-04-25	2
4	100055_20190912	09/12/2019 12:00:00 AM	82	Routine - Unscheduled	100055	2019-09-12	2



## Question 4b

Let's look at the 20 businesses with the lowest **median** score. Order your results by the median score followed by the business id to break ties. The resulting table should look like:

*Hint: You may find the `as_index` argument in the `groupby` method important. [The documentation is linked here!](#)*

	bid		name	median score
<b>3876</b>	84590	Chaat Corner		54.0
<b>4564</b>	90622	Taqueria Lolita		57.0
<b>4990</b>	94351	VBowls LLC		58.0

<b>2719</b>	69282	New Jumbo Seafood Restaurant	60.5
<b>222</b>	1154	SUNFLOWER RESTAURANT	63.5
<b>1991</b>	39776	Duc Loi Supermarket	64.0
<b>2734</b>	69397	Minna SF Group LLC	64.0
<b>3291</b>	78328	Golden Wok	64.0
<b>4870</b>	93150	Chez Beesen	64.0
<b>4911</b>	93502	Smoky Man	64.0
<b>5510</b>	98995	Vallarta's Taco Bar	64.0
<b>1457</b>	10877	CHINA FIRST INC.	64.5
<b>2890</b>	71310	Golden King Vietnamese Restaurant	64.5
<b>4352</b>	89070	Lafayette Coffee Shop	64.5
<b>505</b>	2542	PETER D'S RESTAURANT	65.0
<b>2874</b>	71008	House of Pancakes	65.0
<b>818</b>	3862	IMPERIAL GARDEN SEAFOOD RESTAURANT	66.0
<b>2141</b>	61427	Nick's Foods	66.0
<b>2954</b>	72176	Wolfes Lunch	66.0
<b>4367</b>	89141	Cha Cha Cha on Mission	66.5

```
In [ ]: twenty_lowest_scoring = ins_named.groupby(['bid', 'name'],
...
twenty_lowest_scoring
```

Out[ ]:

	<b>bid</b>	<b>name</b>	<b>score</b>
<b>3876</b>	84590	Chaat Corner	54.0
<b>4564</b>	90622	Taqueria Lolita	57.0
<b>4990</b>	94351	VBowls LLC	58.0
<b>2719</b>	69282	New Jumbo Seafood Restaurant	60.5
<b>222</b>	1154	SUNFLOWER RESTAURANT	63.5
<b>1991</b>	39776	Duc Loi Supermarket	64.0
<b>2734</b>	69397	Minna SF Group LLC	64.0
<b>4870</b>	93150	Chez Beesen	64.0

<b>4911</b>	93502	Smoky Man	64.0
<b>5510</b>	98995	Vallarta's Taco Bar	64.0
<b>1457</b>	10877	CHINA FIRST INC.	64.5
<b>2890</b>	71310	Golden King Vietnamese Restaurant	64.5
<b>4352</b>	89070	Lafayette Coffee Shop	64.5
<b>505</b>	2542	PETER D'S RESTAURANT	65.0
<b>2141</b>	61427	Nick's Foods	66.0
<b>2954</b>	72176	Wolfes Lunch	66.0
<b>4367</b>	89141	Cha Cha Cha on Mission	66.5
<b>4726</b>	91843	Hello Sandwich & Noodle	66.5
<b>2985</b>	74216	MDJK Food Service	67.0
<b>1710</b>	32887	Mission Beach Cafe	67.5

---

## Question 4c

Let's now examine the descriptions of violations for inspections with `score > 0` and `score < 65`. Construct a **Series** indexed by the `description` of the violation from the `vio` table with the value being the number of times that violation occurred for inspections with the above score range. Sort the results in descending order of the count.

The first few entries should look like:

```
Unclean or unsanitary food contact surfaces
43
High risk food holding temperature
42
Unclean or degraded floors walls or ceilings
40
```

## Unapproved or unmaintained equipment or utensils 39

You will need to use `merge` twice.

```
In [ ]: ins_new = ins_named[ins_named['score'] < 65]
low_score_violations = pd.merge(ins_new, ins2vio, on='iid')
low_score_violations = pd.merge(low_score_violations, vio,
low_score_violations = low_score_violations.groupby('description')
low_score_violations.head(20)
```

```
Out[ ]: description
Unclean or unsanitary food contact surfaces
43
High risk food holding temperature
42
Unclean or degraded floors walls or ceilings
40
Unapproved or unmaintained equipment or utensils
39
Foods not protected from contamination
37
High risk vermin infestation
37
Inadequate food safety knowledge or lack of certified food
safety manager      35
Inadequate and inaccessible handwashing facilities
35
Improper thawing methods
30
Unclean hands or improper use of gloves
27
Improper cooling methods
25
Unclean nonfood contact surfaces
21
Inadequately cleaned or sanitized food contact surfaces
20
Improper food storage
20
Contaminated or adulterated food
18
Moderate risk vermin infestation
15
Permit license or inspection report not posted
13
Moderate risk food holding temperature
13
Food safety certificate or food handler card not available
12
```

Improper storage use or identification of toxic substances  
10  
dtype: int64

## Question 4d

Let's figure out which restaurant had the worst scores ever (single lowest score).

**In the cell below, write the name of the restaurant** with the lowest inspection scores ever. You can also head to [yelp.com](https://www.yelp.com) and look up the reviews page for this restaurant. Feel free to add anything interesting you want to share.

```
In [ ]: ins_named.sort_values('score', ascending=True).head()
```

```
Out[ ]:
```

	iid	date	score	type	bid	timestamp
<b>10898</b>	86718_20180522	05/22/2018 12:00:00 AM	45	Routine - Unscheduled	86718	2018-05- 22
<b>291</b>	1154_20190327	03/27/2019 12:00:00 AM	46	Routine - Unscheduled	1154	2019-03- 27
<b>236</b>	10877_20190701	07/01/2019 12:00:00 AM	48	Routine - Unscheduled	10877	2019-07- 01
<b>6433</b>	67237_20180914	09/14/2018 12:00:00 AM	51	Routine - Unscheduled	67237	2018-09- 14
<b>10285</b>	84590_20181001	10/01/2018 12:00:00 AM	54	Routine - Unscheduled	84590	2018-10- 01

```
In [ ]: worst_restaurant = 'Lollipop'  
worst_restaurant
```

```
Out[ ]: 'Lollipop'
```

**Congratulations! You have finished  
Homework 2!**