

LECTURE 16

SQL I

SQL and Databases: An alternative to Pandas and CSV files.

Why Databases

- **Why Databases**
- Intro to SQL
- Tables and Schema
- Basic Queries
- Grouping

Previously: CSV Files and Pandas

we've usually worked with data stored in CSV files.

Calls_for_Service.csv

pd.read_csv

	CASENO	OFFENSE	EVENTDT	EVENTTM	CVLEGEND	CVDOW	InDbDate	Block_Location	BLKADDR	City	State
0	21014296	THEFT MISD. (UNDER \$950)	04/01/2021 12:00:00 AM	10:58	LARCENY	4	06/15/2021 12:00:00 AM	Berkeley, CA\n(37.869058, -122.270455)	NaN	Berkeley	CA
1	21014391	THEFT MISD. (UNDER \$950)	04/01/2021 12:00:00 AM	10:38	LARCENY	4	06/15/2021 12:00:00 AM	Berkeley, CA\n(37.869058, -122.270455)	NaN	Berkeley	CA
2	21090494	THEFT MISD. (UNDER \$950)	04/19/2021 12:00:00 AM	12:15	LARCENY	1	06/15/2021 12:00:00 AM	2100 BLOCK HASTE ST\nBerkeley, CA\n(37.864908,...	2100 BLOCK HASTE ST	Berkeley	CA
3	21090204	THEFT FELONY (OVER \$950)	02/13/2021 12:00:00 AM	17:00	LARCENY	6	06/15/2021 12:00:00 AM	2600 BLOCK WARRING ST\nBerkeley, CA\n(37.86393...)	2600 BLOCK WARRING ST	Berkeley	CA
4	21090179	BURGLARY AUTO	02/08/2021 12:00:00 AM	6:20	BURGLARY - VEHICLE	1	06/15/2021 12:00:00 AM	2700 BLOCK GARBER ST\nBerkeley, CA\n(37.86066,...	2700 BLOCK GARBER ST	Berkeley	CA

Perfectly reasonable workflow for small data that we're not actively sharing with others.



Brief Databases Overview

A **database** is an organized collection of data.

A **Database Management System (DBMS)** is a software system that **stores, manages**, and **facilitates access** to one or more databases.



Advantages of DBMS over CSV (or Similar)

Data Storage:

- **Reliable storage** to survive system crashes and disk failures.
- Optimize to **compute on data that does not fit in memory**.
- Special data structures to **improve performance**.

Data Management:

- Configure how data is **logically organized** and **who has access**.
- Can enforce guarantees on the data (e.g. non-negative person weight or age).
 - Can be used to **prevent data anomalies**.
 - Ensures **safe concurrent operations** on data (multiple users reading and writing simultaneously, e.g. ATM transactions).

Intro to SQL

- Why Databases
- **Intro to SQL**
- Tables and Schema
- Basic Queries
- Grouping

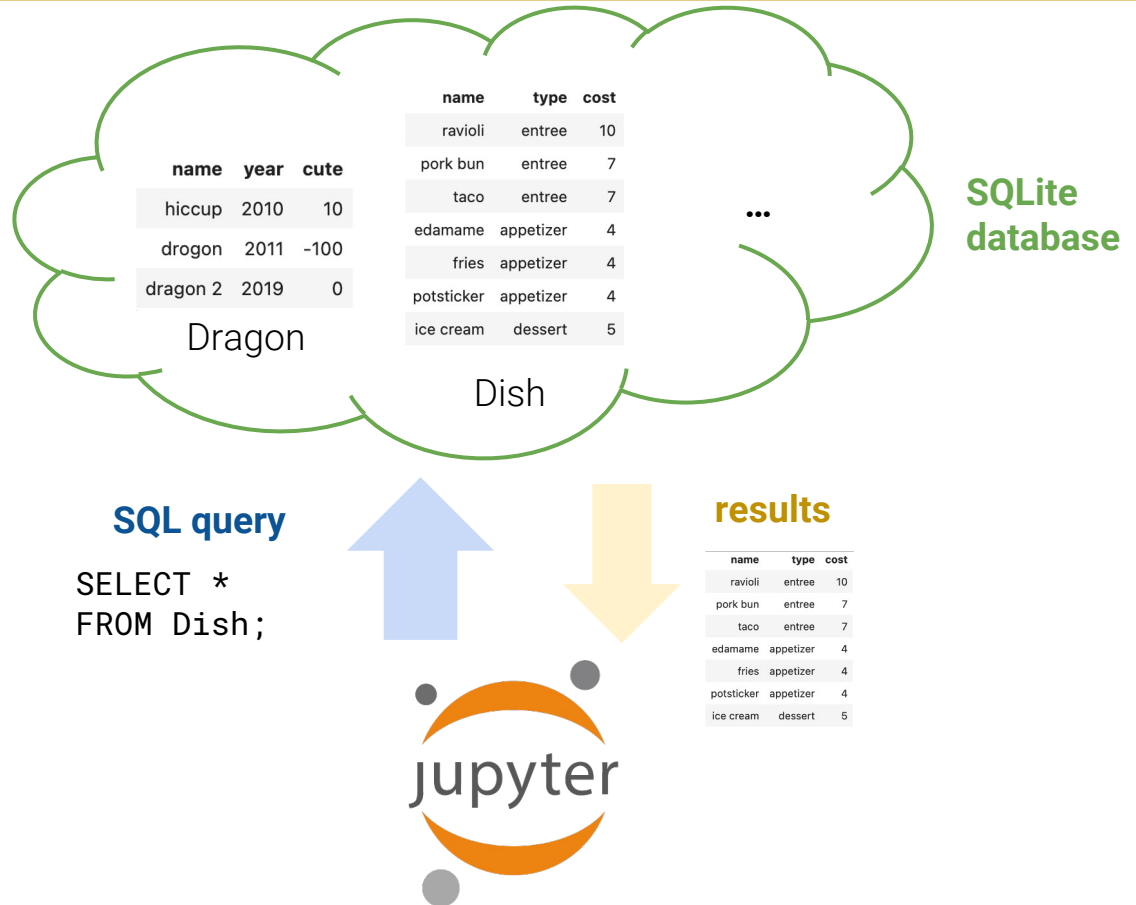
Today we'll be using a programming language called "Structured Query Language" or **SQL**.

- SQL is its own programming language, totally distinct from Python.
- SQL is a special purpose programming language used specifically for communicating with databases.
- We will program in SQL using Jupyter notebooks.

How to pronounce? An ongoing [debate](#).

Let's see a quick demo of how we can use SQL to connect to a database and view a SQL table.

Quick SQL Overview



Demo

Step 1: Load the SQL Module

Our first step is to load the SQL module. We do so using the `ipython cell magic` command:

```
%load_ext sql
```

Step 2: Connect to a Database

Our first step is to load the SQL module. We do so using the `ipython cell magic` command:

```
%load_ext sql
```

The second step is to connect to a database. We use the `%%sql` header to tell Jupyter that this cell represents SQL code rather than Python code.

```
%%sql  
sqlite:///data/basic_examples.db
```

(A note about SQLite)

Our first step is to load the SQL module. We do so using the `ipython cell magic` command:

```
%load_ext sql
```

The second step is to connect to a database.

We use the `%%sql` header to tell Jupyter that this cell represents SQL code rather than Python code.

```
%%sql  
sqlite:///data/basic_examples.db
```

Connected: @data/18_basic_examples.db

In real world practice, you'd probably connect to a remote server.

There are various extensions to SQL.

We are learning the SQL commands and syntax supported by the SQLite library.

If you're curious: **SQLite** is a library that provides a relational DBMS (RDBMS). It is lightweight and offers file-based databases.

Tables and Schema

- Why Databases
- Intro to SQL
- **Tables and Schema**
- Basic Queries
- Grouping

Column or Attribute or Field

Row or
Record or
Tuple

name TEXT, PK	year INT, >=2000	cute INT
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon

← table name

SQL tables are also called relations.
SQL Style: Use *singular, CamelCase* names for SQL tables! For more, see [this post](#).



Column or Attribute or Field

Row or Record or Tuple

name TEXT, PK	year INT, >=2000	cute INT
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Column Properties
ColName,
Type, **Constraint**

Dragon ← table name

SQL **tables** are also called **relations**.

SQL Style: Use *singular, CamelCase* names for SQL tables! For more, see [this post](#).

Every column in a SQL table has three properties: **ColName**, **Type**, and zero or more **Constraints**.
(Contrast with **pandas: Series** have names and types, but no constraints.)

Table Schema

A **schema** describes the logical structure of a table. Whenever a new table is created, the creator must declare its schema.

For each column, specify the:

- **Column name**
- **Data type**
- **Constraint(s) on values**

```
CREATE TABLE Dragon (  
  name TEXT PRIMARY KEY,  
  year INTEGER CHECK (year >= 2000),  
  cute INTEGER  
)
```

Repeat for all tables in the database (see demo nb):

type	name	tbl_name	rootpage	sql
table	sqlite_sequence	sqlite_sequence	7	CREATE TABLE sqlite_sequence(name,seq)
table	Dragon	Dragon	2	CREATE TABLE Dragon (name TEXT PRIMARY KEY, year INTEGER CHECK (year >= 2000), cute INTEGER)
table	Dish	Dish	4	CREATE TABLE Dish (name TEXT PRIMARY KEY, type TEXT, cost INTEGER CHECK (cost >= 0))

Example Types

Some examples of SQL **types**:

- **INT**: Integers.
- **FLOAT**: Floating point numbers.
- **TEXT**: Strings of text.
- **BLOB**: Arbitrary data, e.g. songs, video files, etc.
- **DATETIME**: A date and time.

Note: Different implementations of SQL support different types.

- SQLite: <https://www.sqlite.org/datatype3.html>
- MySQL: <https://dev.mysql.com/doc/refman/8.0/en/data-types.html>

we will use SQLite!

Example Constraints

Some examples of **constraints**:

- **CHECK**: data must obey the given check constraint.
- **PRIMARY KEY**: specifies that this key is used to uniquely identify rows in the table.
- **NOT NULL**: null data cannot be inserted for this column.
- **DEFAULT**: provides a default value to use if user does not specify on insertion.

type	name	tbl_name	rootpage	sql
table	sqlite_sequence	sqlite_sequence	7	CREATE TABLE sqlite_sequence(name,seq)
table	Dragon	Dragon	2	CREATE TABLE Dragon (name TEXT PRIMARY KEY, year INTEGER CHECK (year >= 2000), cute INTEGER)
table	Dish	Dish	4	CREATE TABLE Dish (name TEXT PRIMARY KEY, type TEXT, cost INTEGER CHECK (cost >= 0))
table	Scene	Scene	6	CREATE TABLE Scene (id INTEGER PRIMARY KEY AUTOINCREMENT, biome TEXT NOT NULL, city TEXT NOT NULL, visitors INTEGER CHECK (visitors >= 0), created_at DATETIME DEFAULT (DATETIME('now')))

What is this
primary key
constraint?

Primary Keys

A **primary key** is the set of column(s) used to uniquely identify each record in the table.

- In the Dragon table, the “**name**” of each Dragon is the primary key.
- In other words, no two dragons can have the same name!
- Primary key is used **under the hood** for all sorts of optimizations.

name TEXT, PK	year INT, >=2000	cute INT
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Why specify primary keys?
More next time when we
discuss JOINS...

Basic Queries

- Why Databases
- Intro to SQL
- Tables and Schema
- **Basic Queries**
- Grouping

```
SELECT <column list>  
FROM <table>
```



**Summary
so far**

;



Marks the end of a SQL statement.



Goal of this section

```
SELECT <column list>  
FROM <table>  
[WHERE <predicate>]  
[ORDER BY <column list>]  
[LIMIT <number of rows>]  
[OFFSET <number of rows>];
```

By the end of this section, you will learn these new keywords!

But first, more SELECT

Recall our simplest query, which returns the full relation:

```
SELECT *  
FROM Dragon;
```



table name

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0
puff	2010	100
smaug	2011	None

SELECT specifies the column(s) that we wish to appear in the output. **FROM** specifies the database table from which to select data.

Every query must include a **SELECT** clause (how else would we know what to return?) and a **FROM** clause (how else would we know where to get the data?)

An asterisk (*) is shorthand for “all columns”. *Let's see a bit more in our demo.*

But first, more SELECT

Recall our simplest query, which returns the full relation:

```
SELECT *  
FROM Dragon;
```



table name

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0
puff	2010	100
smaug	2011	None

We can also SELECT only a **subset of the columns**:

```
SELECT cute, year  
FROM Dragon;
```

column expression list

cute	year
10	2010
-100	2011
0	2019
100	2010
None	2011



Columns selected
in specified order

Aliasing with AS

To rename a **SELECT**ed column, use the **AS** keyword

```
SELECT cute AS cuteness,  
       year AS birth  
FROM Dragon;
```

An **alias** is a name given to a column or table by a programmer. Here, “cuteness” is an alias of the original “cute” column (and “birth” is an alias of “year”)

cuteness	birth
10	2010
-100	2011
0	2019
100	2010
None	2011

SQL Style: Newline Separators

The following two queries both retrieve the same relation:

```
SELECT cute AS cuteness,  
       year AS birth  
FROM Dragon;
```

(more readable)



cuteness	birth
10	2010
-100	2011
0	2019
100	2010
None	2011

```
SELECT cute AS  
       cuteness, year AS  
       birth FROM Dragon;
```



Use newlines and whitespace wisely in your SQL queries. It will simplify your debugging process!


Uniqueness with DISTINCT

To return only unique values, combine **SELECT** with the **DISTINCT** keyword

```
SELECT DISTINCT year
FROM Dragon;
```

Notice that 2010 and 2011 only appear once each in the output.

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0
puff	2010	100
smaug	2011	None



year
2010
2011
2019



WHERE: Select a rows based on conditions

To select only some rows of a table, we can use the **WHERE** keyword.

```
SELECT name, year
FROM Dragon
WHERE cute > 0;
```

condition

name	year
hiccup	2010
puff	2010

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0
puff	2010	100
smaug	2011	None

Dragon

WHERE: Select a rows based on conditions

Comparators **OR**, **AND**, and **NOT** let us form more complex conditions.

```
SELECT name, year
FROM Dragon
WHERE cute > 0 OR year > 2013;
      condition
```

name	cute	year
hiccup	10	2010
dragon 2	0	2019
puff	100	2010

Check if values are contained IN a specified list

```
SELECT name, year
FROM Dragon
WHERE name IN ('puff', 'hiccup');
```

name	year
hiccup	2010
puff	2010

WHERE with NULL Values

NULL (the SQL equivalent of NaN) is stored in a special format – we can't use the "standard" operators =, >, and <.

Instead, check if something **IS** or **IS NOT NULL**

```
SELECT name, year
FROM Dragon
WHERE year IS NOT NULL;
```

name	cute
hiccup	10
drogon	-100
dragon 2	0
puff	100

Always work with NULLs using the **IS** operator.
NULL cannot work with standard comparisons:
in fact, NULL = NULL actually returns False!

Specify which column(s) we should order the data by

```
SELECT *  
FROM Dragon  
ORDER BY cute DESC;
```


column



(by default, SQL orders by
ascending order: **ASC**)

name	year	cute
puff	2010	100
hiccup	2010	10
dragon 2	2019	0
drogon	2011	-100
smaug	2011	None

ORDER BY: Sort rows

Specify which column(s) we should order the data by

```
SELECT *  
FROM Dragon  
ORDER BY year, cute DESC;
```

Can also order by multiple
columns (for tiebreaks)

Sorts **year** in ascending order and **cute** in descending order. If you want **year** to be ordered in descending order as well, you need to specify **year DESC, cute DESC**;

name	year	cute
puff	2010	100
hiccup	2010	10
drogon	2011	-100
smaug	2011	None
dragon 2	2019	0

OFFSET and LIMIT?

1. SELECT *
FROM Dragon
LIMIT 2;

A.

name	year	cute
hiccup	2010	10
drogon	2011	-100

2. SELECT *
FROM Dragon
LIMIT 2
OFFSET 1;

B.

name	year	cute
drogon	2011	-100
dragon 2	2019	0

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon

Matching: Which query matches each relation?

What do you think the **LIMIT** and **OFFSET** keywords do?

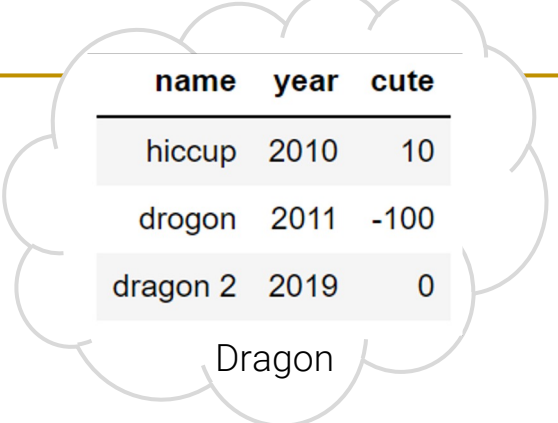


OFFSET and LIMIT

The **LIMIT** keyword lets you retrieve N rows (like **pandas head**).

```
SELECT *  
FROM Dragon  
LIMIT 2;
```

name	year	cute
hiccup	2010	10
drogon	2011	-100




name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon

The **OFFSET** keyword tells SQL to skip the first N rows of the output, then apply **LIMIT**.

```
SELECT *  
FROM Dragon  
LIMIT 2  
OFFSET 1;
```

name	year	cute
drogon	2011	-100
dragon 2	2019	0

 Unless you use **ORDER BY**, there is **no guaranteed order** of rows in the relation!



Summary so far

```
SELECT <column list>  
FROM <table>  
[WHERE <predicate>]  
[ORDER BY <column list>]  
[LIMIT <number of rows>]  
[OFFSET <number of rows>];
```

- All queries must include **SELECT** and **FROM**. The remaining keywords are optional.
- By convention, use **all caps** for keywords in SQL statements.
- Use **newlines** to make code more readable.

Grouping

- Why Databases
- Intro to SQL
- Tables and Schema
- Basic Queries
- **Grouping**

The Dish Table

We're ready for a more complicated table.

```
SELECT *  
FROM Dish;
```

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

The Dish Table

We're ready for a more complicated table.

```
SELECT *  
FROM Dish;
```

Notice the repeated dish [types](#). What if we wanted to investigate trends across each group?

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Order of operations: SELECT → FROM → WHERE → GROUP BY

```
SELECT type, SUM(cost)
FROM Dish
GROUP BY type;
```

Correct! 

```
GROUP BY type
SELECT type, SUM(cost)
FROM Dish;
```

Incorrect 

Always follow the SQL order of operations. Let SQL take care of the rest.

GROUP BY

GROUP BY is similar to pandas `groupby()`.

```
SELECT type  
FROM Dragon  
GROUP BY type;
```

type		name	type	cost
appetizer	entree	ravioli	entree	10
		ramen	entree	13
		taco	entree	7
dessert	appetizer	edamame	appetizer	4
		fries	appetizer	4
		potsticker	appetizer	4
entree	dessert	ice cream	dessert	5

Aggregating Across Groups

Like pandas, SQL has **aggregate functions**: MAX, SUM, AVG, FIRST, etc.

For more aggregations, see: https://www.sqlite.org/lang_aggfunc.html

```
SELECT type, SUM(cost)
FROM Dish
GROUP BY type;
```

type	SUM(cost)
appetizer	12
dessert	5
entree	30

Wait, something's weird...

Wait, something's weird...

```
SELECT type, SUM(cost)
FROM Dish
GROUP BY type;
```

We told SQL to SUM in our SELECT statement...

...but didn't specify the groups until GROUP BY

This is okay!

Unlike Python, SQL is a **declarative programming language**.

Declarative programming is a non-imperative style of programming in which programs describe their desired results without explicitly listing commands or steps that must be performed.

[Wikipedia](#)

Declarative programming is a non-imperative style of programming in which programs describe their desired results without explicitly listing commands or steps that must be performed.

[Wikipedia](#)

What this means to us:

- We “declare” our desired end result
- SQL handles the rest! We do *not* need to specify any logical steps for how this result should be created

We just need to follow the **SQL order of operations** with our clauses to allow SQL to parse our request. Everything else will be handled behind the scenes.

High-level cheat sheet on order of **execution** by the SQL engine ([more info](#)):

- | | |
|----------|-------------|
| 1. FROM | 4. GROUP BY |
| 2. JOIN | 5. SELECT |
| 3. WHERE | 6. ORDER BY |

Using Multiple Aggregation Functions

```
SELECT type,  
       SUM(cost),  
       MIN(cost),  
       MAX(name)  
FROM Dish  
GROUP BY type;
```

What do you think will happen?

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



Using Multiple Aggregation Functions

```
SELECT type,  
       SUM(cost),  
       MIN(cost),  
       MAX(name)  
FROM Dish  
GROUP BY type;
```



type	SUM(cost)	MIN(cost)	MAX(name)
appetizer	12	4	potsticker
dessert	5	5	ice cream
entree	30	7	taco

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

This was much more difficult in pandas!



The COUNT Aggregation

COUNT is used to count the number of rows belonging to a group.

```
SELECT year, COUNT(cute)
FROM Dragon
GROUP BY year;
```

Similar to pandas `groupby().count()`

year	COUNT(cute)
2010	2
2011	1
2019	1

```
SELECT year, COUNT(*)
FROM Dragon
GROUP BY year;
```

Similar to pandas
`groupby().size()`

year	COUNT(*)
2010	2
2011	2
2019	1

COUNT(*) returns the number of rows in each group, including rows with **NULLs**.



Summary so far

```
SELECT <column expression list>  
FROM <table>  
[WHERE <predicate>]  
[GROUP BY <column list>]  
[ORDER BY <column list>]  
[LIMIT <number of rows>]  
[OFFSET <number of rows>];
```

- By convention, use **all caps** for keywords in SQL statements.
- Use **newlines** to make SQL code more readable.
- **AS** keyword: rename columns during selection process.
- **Column Expressions may include aggregation functions (MAX, MIN, etc.)**